

Exact Paradoxical Collatz Cylinders: A Formal Search Note

Collatz proof project*

September 5, 2026

Abstract

We formalize an exact quotient-interval characterization of paradoxical shortcut Collatz segments. It permits exhaustive enumeration at each fixed length without a seed-height cutoff. Lean proves the complete length-eight classification above seed two. A separate exact-integer Python enumeration through length twenty reproduces the same five segments. A suffix-correction bound now prunes whole prefix subtrees. Tested searches at selected lengths through forty complete with substantially fewer nodes than full residue enumeration. This remains a fixed-length research instrument, not a proof of global finiteness or Collatz.

1 Definition and quotient interval

Let $T(n) = n/2$ for even n and $(3n + 1)/2$ for odd n . Let $j_L(n)$ count odd steps before time L . Following Rozier and Terracol [1], a segment is paradoxical when

$$n > 0, \quad T^L(n) \geq n, \quad 3^{j_L(n)} < 2^L.$$

Intermediate descent is allowed. For a fixed residue $0 \leq r < 2^L$, set $u = T^L(r)$ and $A = 3^{j_L(r)}$. The affine transport identity is

$$T^L(r + 2^L q) = u + Aq,$$

and the odd-step count is constant on this residue class.

Theorem 1. *Assume $A < 2^L$. For every natural q , the seed $r + 2^L q$ is paradoxical exactly when*

$$0 < r + 2^L q, \quad r \leq u, \quad q \leq \left\lfloor \frac{u - r}{2^L - A} \right\rfloor.$$

Proof. Non-descent is equivalent to $r + 2^L q \leq u + Aq$, or $r + (2^L - A)q \leq u$. Since the coefficient is positive, this is precisely the stated interval and $r \leq u$. Positivity excludes the zero seed. In Lean, retaining $r \leq u$ is essential because natural subtraction truncates at zero. $\square$

To enumerate seeds at least three, impose also $q \geq \max(0, \lceil (3 - r)/2^L \rceil)$. Noncontracting residues contribute none. The unique decomposition $n = (n \bmod 2^L) + 2^L \lfloor n/2^L \rfloor$ proves coverage of all seeds at the selected length.

*Prepared with Codex (OpenAI). This is a formalization and small-case reproduction, not a literature-priority claim.

2 Checked small case and experiment

Theorem 2. *For every $n > 2$, the length-eight segment is paradoxical if and only if*

$$n \in \{7, 9, 18, 19, 25\}.$$

Proof. A kernel-decided table checks all 256 residues. Every contracting residue with $u \geq r$ satisfies $u-r < 256-A$, so its admissible quotient is zero. The possible residues are 0, 1, 2, 7, 9, 18, 19, 25; the first three are excluded by $n > 2$. The five remaining trajectories directly satisfy the definition. $\square$

Seed	Endpoint at eight	Odd steps	First descent time
7	8	5	7
9	10	5	2
18	20	5	1
19	20	5	4
25	26	5	2

The Python census exhausts lengths one through twenty, using all 2^L residues at each length and the full quotient interval. It finds only these five pairs. Each reported segment is replayed directly. Unit tests independently check every residue through length ten and interval membership and boundaries through length eight. The length-twenty census is a tested Python computation, not a Lean-extracted or kernel-certified enumeration. Only the displayed length-eight classification has been proved in Lean.

3 Research role and reproducibility

Rozier and Terracol prove that finiteness of paradoxical segments above seed two would imply Collatz [1]. Their reduction is already represented in the nested source `rwst__lean-code/RT/CRozLemma32.lean`. Its local toolchain differs from the main project's; that nested headline theorem was source-inspected, not rebuilt or axiom-audited in this pass. It is not imported by the new module. Other nearby library results use cited computational assumptions, which must not be silently imported as proofs.

The new main module is `lean/Collatz/Paradoxical.lean`. Its principal declarations are

```
paradoxical_quotient_interval
paradoxical_residue_iff
paradoxical_eight_iff
```

These are included in the expanded 64-declaration audit, with only standard foundational axioms. No new axiom, `sorry`, or `native_decide` is used. The table uses ordinary kernel-checked `decide`.

Build with `lake build Collatz.Paradoxical` from `lean/`. From the repository root, run `python3 analysis/paradoxical_cylinders.py`; the JSON output records the exact finite scope. Brute-force residue enumeration costs exponentially many states. The prefix pruning extension below avoids many of those states, without establishing all-length termination. The missing task is an independently proved restriction on the contracting words that could support arbitrarily long paradoxical segments. Finite experimental absence does not supply that restriction.

4 A proved prefix-pruning rule

Represent a parity word by Boolean letters in orbit order. For a word w , let $|w|$ be its length, $j(w)$ its number of ones, and $d(w)$ its affine correction. The exact composition identity is

$$d(uv) = 3^{j(v)}d(u) + 2^{|u|}d(v).$$

The maximal suffix correction at fixed length h and odd count k is bounded by

$$d(v) \leq 2^{h-k}(3^k - 2^k). \quad (1)$$

This is the same extremal correction bound appearing in Rozier–Terracol [1], Theorem 2.2, translated from the remainder $d(v)/2^h$ to the integer correction. The new module proves it directly by induction, without importing that paper’s formalization.

Theorem 3. *Let a prefix u have length s , odd count a , and correction D . Every suffix v of length h and odd count k satisfies*

$$d(uv) \leq U_k := 3^k D + 2^{s+h-k}(3^k - 2^k).$$

Suppose a seed $n \geq m$ realizes uv and $3^{a+k} < 2^{s+h}$. If

$$U_k < (2^{s+h} - 3^{a+k})m, \quad (2)$$

then its segment of length $s + h$ is not paradoxical.

Proof. The upper bound follows from composition and (1). If the endpoint is at least n , the exact affine identity gives

$$(2^{s+h} - 3^{a+k})n \leq d(uv) \leq U_k.$$

The positive coefficient and $n \geq m$ contradict (2). □

For a canonical prefix residue $r \bmod 2^s$, the smallest seed at least three in the class is

$$m = r + 2^s \max\left(0, \left\lceil \frac{3 - r}{2^s} \right\rceil\right).$$

Check all $k = 0, \dots, h$. A noncontracting count contributes no paradoxical completion. If every contracting count passes the strict pruning test, the entire prefix subtree can be discarded. No monotonicity shortcut over the possible counts is assumed in the implementation.

The induction behind (1) first proves the division-free inequality

$$2^{j(w)}(d(w) + 2^{|w|}) \leq 2^{|w|}3^{j(w)}.$$

Then $j(w) \leq |w|$ and cancellation of $2^{j(w)} > 0$ give (1). These steps and the exclusion theorem are checked in `Collatz.ParadoxicalPruning`.

5 Search results and verification boundary

The implementation `analysis/paradoxical_pruned.py` follows residue lifts in a depth-first tree. Each node stores its depth, residue, endpoint, odd count, and correction. The affine identity is checked at every visited node. Surviving leaves use the exact quotient interval from Section 1, so there is no seed-height cutoff.

Length	Visited nodes	Full residue count	Segments	Status
8	201	256	5	Complete
20	1,095	1,048,576	0	Complete
27	38,425	134,217,728	50	Complete
32	28,443	4,294,967,296	0	Complete
40	123,887	1,099,511,627,776	0	Complete
65	2,000,000	2^{65}	—	Incomplete

These are selected lengths, not every length through forty. Tree nodes and full-scan residues are different units; both counts are provided explicitly. The length-65 run reaches its work limit with 21 pending nodes. Its lack of emitted segments does not establish absence of segments.

The six Python controls include exhaustive comparisons through length fourteen, all short suffix bounds through length eight, direct residue and endpoint checks, and a test that a work-limited run cannot report completion. Every emitted segment is independently replayed with the shortcut rule. The JSON file `analysis/paradoxical_pruned_results.json` preserves the status and discovered segments.

Formal and computational claims

The added module proves `correction_upper`, the completion bound, and the pruning implication. These mathematical statements are in Lean. The Python tree traversal and the reported larger-length counts are tested computations, not Lean-extracted code or kernel-verified search certificates. The existing all-seed length-eight classification remains separately kernel-proved. The selected 64-declaration audit includes the new bounds and reports only standard foundational axioms.

Build from `lean/` with

```
lake build Collatz.ParadoxicalPruning
```

From the repository root, reproduce the experiment with

```
python3 analysis/paradoxical_pruned.py
```

The result makes fixed-length investigation substantially more practical, but does not yield a uniform cutoff on possible lengths. The length-65 work limit is an observed computational boundary, not a mathematical obstruction. A restriction uniform in length or circuit count is still needed to obtain global paradoxical finiteness and use the existing reduction to Collatz.

References

- [1] O. Rozier and C. Terracol, *Paradoxical behavior in Collatz sequences*, arXiv:2502.00948v3 (2025), Definition 1.1 and Section 3. <https://arxiv.org/html/2502.00948v3>.